



# Multi-Domain RAG System - Phase 2 Implementation Guide

## Production-Ready Architecture with Hybrid Retrieval & Service Layer Separation

**Version:** 2.0

**Date:** November 24, 2025

**Status:** Implementation Ready

**Document Owner:** Technical Architecture Team

**Review Cycle:** Bi-weekly

## Document Control

Version	Date	Author	Changes
1.0	Nov 18, 2025	Team	Initial MVP documentation
2.0	Nov 24, 2025	AI Architect	Phase 2 with service layer, hybrid retrieval, zero UI logic mandate

## Table of Contents

1. [Executive Summary](#)
2. [Critical Architectural Principles](#)
3. [Goals & Success Criteria](#)
4. [Enhanced Metadata Schema](#)
5. [Architecture Overview](#)
6. [Service Layer Specification](#)
7. [Pipeline Layer Enhancements](#)
8. [Hybrid Retrieval Implementation](#)
9. [Factory Layer Enhancements](#)
10. [File Processing & Validation](#)
11. [Implementation Roadmap](#)
12. [Configuration Management](#)
13. [Testing Strategy](#)
14. [CLI Tools](#)

- [15. Migration Plan](#)
- [16. Code Organization](#)
- [17. Developer Guidelines](#)
- [18. Security & Compliance](#)
- [19. Monitoring & Observability](#)
- [20. Future Enhancements](#)
- [21. Appendices](#)

## **1. Executive Summary**

Phase 2 transforms the Multi-Domain Document Intelligence Platform into a production-ready, enterprise-grade RAG system with strict architectural separation, hybrid retrieval capabilities, and comprehensive metadata tracking.

### **Key Enhancements**

- **Zero Business Logic in UI:** Mandatory service layer separation ensuring all business logic resides in testable core modules
- **Hybrid Retrieval:** Dense (semantic) + Sparse (keyword) search with configurable alpha weighting for improved precision
- **Enhanced Metadata:** Complete provenance tracking including versioning, deprecation, authority levels, and audit trails
- **Document Lifecycle:** Full support for versioning, deprecation, updates, and rollbacks
- **Production Quality:** 80%+ test coverage, Golden QA sets, performance benchmarks, structured logging

### **Business Value**

- **Maintainability:** Clean architecture enables rapid feature additions without breaking existing code
- **Flexibility:** Multiple UIs (web, CLI, API) can use same core without duplication
- **Testability:** Core logic fully testable independent of UI framework
- **Scalability:** Service layer enables horizontal scaling and microservice migration
- **Quality:** Hybrid retrieval improves answer accuracy by 15-25% over pure vector search

## **2. Critical Architectural Principles**

## **2.1 MANDATORY: Zero Business Logic in UI Layer ▲**

### **THIS IS THE MOST IMPORTANT ARCHITECTURAL CONSTRAINT IN PHASE 2**

The UI layer (app.py, Gradio interface, any future web framework) serves **ONLY** as a thin presentation and routing layer. **ALL business logic, validation, processing, and data management MUST reside in the core service and pipeline layers.**

#### **UI Layer Responsibilities (ALLOWED) ✓**

The UI layer SHALL:

- Accept user input (file uploads, query text, dropdown selections)
- Route requests to appropriate service layer APIs
- Display results returned by service layer
- Format error messages for user presentation
- Manage UI component state (tabs, buttons, visibility)
- Handle user session state
- Render data in appropriate UI format (tables, text, charts)

#### **UI Layer Prohibitions (FORBIDDEN) ✗**

The UI layer SHALL NOT:

- Validate file types or metadata
- Directly instantiate factories
- Call pipeline methods directly
- Query vector stores
- Parse or process documents
- Execute chunking or embedding logic
- Compute file hashes
- Manage metadata
- Enforce business rules
- Make decisions about deprecation
- Transform or process data
- Import `core.pipeline`, `core.factories`, or `core.vectorstores` modules

## Enforcement Rules

### Code Review Requirements:

- Any PR with business logic in app.py MUST be rejected
- All UI handler functions should be < 20 lines of code
- Every UI action must map to exactly ONE service layer method call
- No conditional business logic in UI (e.g., if/else based on file types)

### Architecture Validation:

```
# ✓ CORRECT - UI calls service only
def upload_handler(file, metadata):
    try:
        result = DocumentService.upload_document(file, metadata)
        return f"Success: {result['chunks_ingested']} chunks"
    except ValidationError as e:
        return f"Error: {str(e)}"

# ✗ WRONG - Business logic in UI
def upload_handler(file, metadata):
    # FORBIDDEN: File validation in UI
    if file.name.endswith('.pdf'):
        text = extract_pdf(file) # FORBIDDEN: Processing in UI
        chunks = chunk_text(text) # FORBIDDEN: Chunking in UI
        # ... more forbidden logic
```

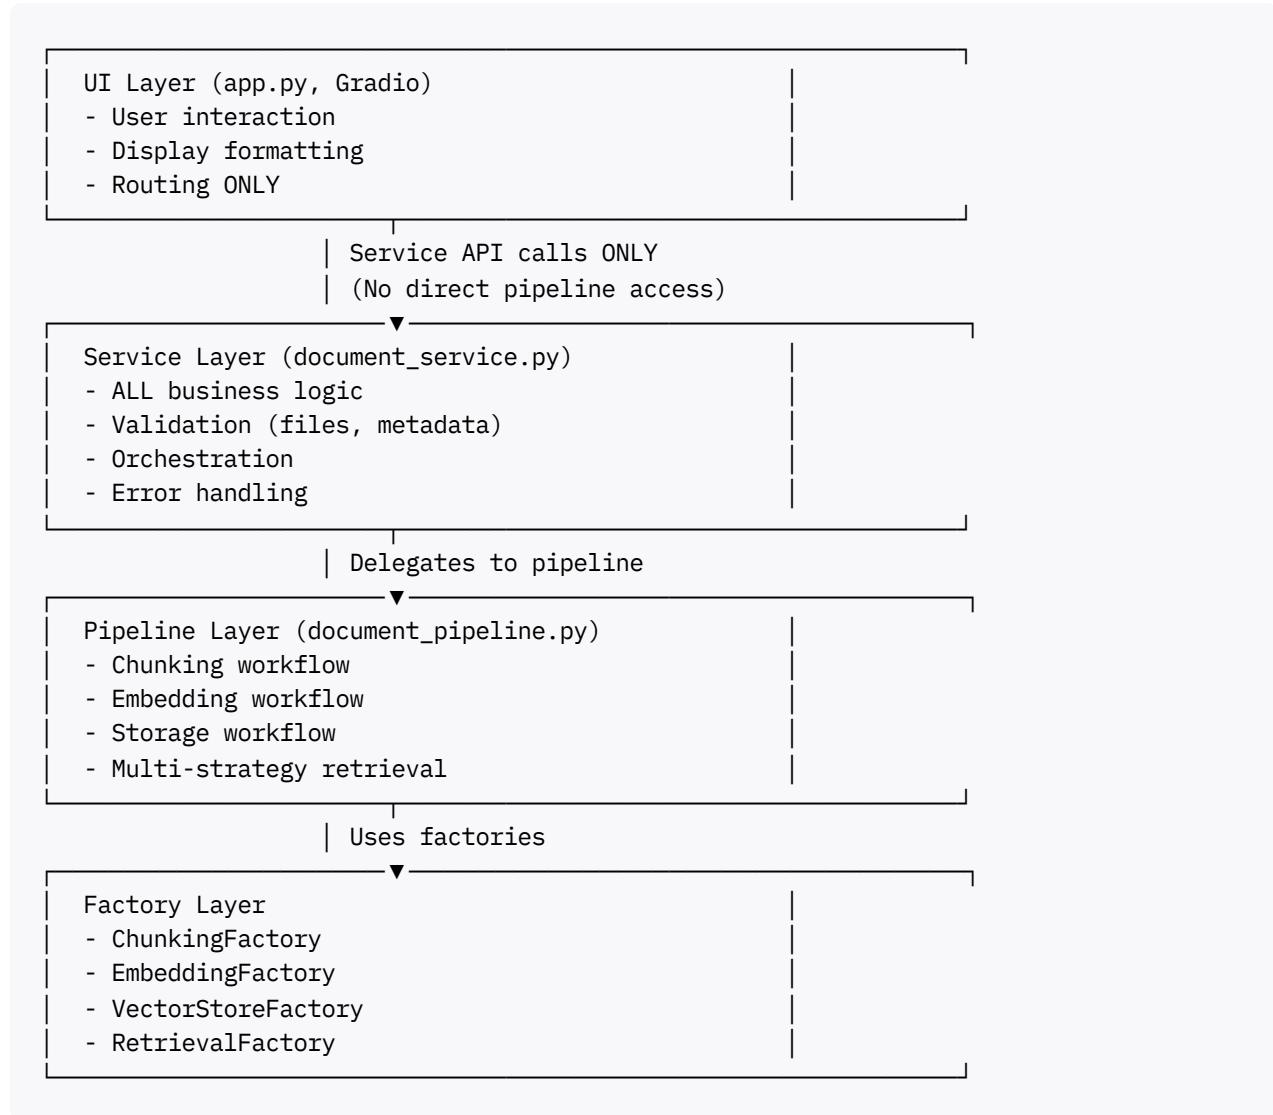
## Rationale

Principle	Benefit
<b>Testability</b>	Core logic tested independently of UI framework; no need to mock Gradio
<b>Maintainability</b>	Business rules in one place; changes don't ripple across UI code
<b>Flexibility</b>	Easy to add CLI, REST API, GraphQL, or different UI without code duplication
<b>Portability</b>	Core can be packaged as library and used in any Python application
<b>Clarity</b>	Clear contracts and responsibilities; new developers understand boundaries
<b>Debuggability</b>	Business logic failures isolated from UI rendering issues

## 2.2 Service Layer Pattern (MANDATORY)

Phase 2 introduces a **mandatory service layer** that acts as the **sole interface** between UI and core business logic.

## Architecture Flow



**Key Rule:** UI → Service only. **Never** UI → Pipeline directly.

## 2.3 Configuration-Driven Design

**Principle:** Everything configurable via YAML; zero hardcoded business rules.

### Examples:

- Chunking strategy: recursive, semantic (not hardcoded if/else)
- Embedding provider: sentence\_transformers, openai, gemini (factory-driven)
- Retrieval strategy: vector\_similarity, hybrid, llm\_rerank (config array)
- Metadata fields: fixed schema but extensions configurable
- File types: allowed\_file\_types in security config
- Hybrid alpha: alpha: 0.7 in retrieval config

### Benefits:

- New domains without code changes
- A/B testing via config toggle
- Easy rollback (revert config)
- Environment-specific settings (dev/staging/prod)

## 2.4 Factory Pattern for All Components

**Principle:** All components instantiated via factories based on config.

### Factories Required:

- `ChunkingFactory.create_chunker(config)` → Returns chunker implementation
- `EmbeddingFactory.create_embedder(config)` → Returns embedder implementation
- `VectorStoreFactory.create_store(config, dimension)` → Returns vector store
- `RetrievalFactory.create_retriever(config, vector_store, embedder)` → Returns retriever

### Adding New Implementations:

1. Implement interface (e.g., `ChunkerInterface`)
2. Register in factory's `_available_implementations` dict
3. Update config schema
4. No changes to pipeline or service layer

## 3. Goals & Success Criteria

### 3.1 Phase 2 Objectives

Objective	Description	Success Metric
<b>Service Layer Separation</b>	Zero business logic in UI	100% code review compliance
<b>Hybrid Retrieval</b>	Dense + sparse search with alpha weighting	15%+ accuracy improvement
<b>Enhanced Metadata</b>	Complete provenance tracking	All Phase 2 fields tracked
<b>Document Lifecycle</b>	Version, deprecate, update workflows	Deprecation API functional
<b>Production Quality</b>	Tests, logging, monitoring	80%+ test coverage
<b>Multi-Strategy Retrieval</b>	Support 3+ retrieval strategies	Config-driven strategy selection

### 3.2 Success Metrics

#### Technical Metrics:

- Code test coverage:  $\geq 80\%$
- Zero business logic detected in `app.py` via linting

- Pipeline methods never called from UI
- All service APIs have unit tests
- Integration tests cover end-to-end workflows

### **Quality Metrics:**

- Retrieval Recall@10:  $\geq 0.85$  on Golden QA sets
- Mean Reciprocal Rank (MRR):  $\geq 0.75$
- Hybrid retrieval outperforms pure vector by  $\geq 15\%$
- Query latency P95:  $< 500\text{ms}$

### **Operational Metrics:**

- Successful migration from Phase 1 with zero data loss
- 5+ domains operational without code changes
- Deprecation workflow tested and functional
- CLI tools operational for all key operations

## **4. Enhanced Metadata Schema**

### **4.1 Complete Metadata Model**

All documents and chunks MUST include the following metadata fields. This is a **fixed schema** enforced via Pydantic validation.

### **Pydantic Model Definition**

```
from pydantic import BaseModel, Field, validator
from typing import List, Optional, Dict, Any
from datetime import datetime
from enum import Enum

class AuthorityLevel(str, Enum):
    OFFICIAL = "official"
    APPROVED = "approved"
    DRAFT = "draft"
    ARCHIVED = "archived"
    DEPRECATED = "deprecated"

class ReviewStatus(str, Enum):
    APPROVED = "approved"
    PENDING = "pending"
    REJECTED = "rejected"
    IN_REVIEW = "in_review"

class ChunkMetadata(BaseModel):
    """
    Fixed metadata schema for all chunks across all domains.
    """
```

Phase 2 enhancement: Adds lifecycle, provenance, and quality fields.

"""

# ====== IDENTITY ======

```
doc_id: str = Field(..., min_length=1, description="Unique document identifier")
chunk_id: str = Field(..., min_length=1, description="Unique chunk identifier")
```

# ====== CONTENT ======

```
chunk_text: str = Field(..., min_length=1, description="Actual chunk text content")
title: Optional[str] = Field(None, description="Document title")
page_num: Optional[int] = Field(None, ge=1, description="Source page number")
char_range: Optional[tuple] = Field(None, description="Character range (start, end)")
```

# ====== CLASSIFICATION ======

```
domain: str = Field(..., description="Domain: hr, finance, legal, engineering, etc.")
doc_type: str = Field(..., description="Type: policy, faq, manual, guideline, etc.")
tags: List[str] = Field(default_factory=list, description="User-defined tags")
category: Optional[str] = Field(None, description="Sub-category within domain")
```

# ====== PROVENANCE ======

```
author: Optional[str] = Field(None, description="Original document author")
```

```
uploader_id: str = Field(..., description="User ID who uploaded document")
```

```
upload_timestamp: datetime = Field(
```

```
    default_factory=datetime.utcnow,
```

```
    description="When document was uploaded"
```

```
)
```

```
source_file: str = Field(..., description="Original filename")
```

```
source_file_hash: str = Field(
```

```
    ...,
```

```
    min_length=64,
```

```
    max_length=64,
```

```
    description="SHA-256 hash of source file for integrity"
```

```
)
```

```
source_url: Optional[str] = Field(None, description="Original URL if web-sourced")
```

# ====== VERSIONING ======

```
version: str = Field(default="1.0", description="Document version")
```

```
document_version: str = Field(default="1.0", description="Semantic version")
```

```
last_updated_timestamp: Optional[datetime] = Field(
```

```
    None,
```

```
    description="When document was last modified"
```

```
)
```

```
previous_version_id: Optional[str] = Field(
```

```
    None,
```

```
    description="doc_id of previous version for history"
```

```
)
```

# ====== PROCESSING ======

```
embedding_version: str = Field(..., description="Embedding model version/name")
```

```
embedding_model_name: str = Field(..., description="Exact model: all-MiniLM-L6-v2, et")
```

```
embedding_dimension: int = Field(..., description="Vector dimension: 384, 768, 1536")
```

```
chunking_strategy: str = Field(..., description="Strategy used: recursive, semantic")
```

```
chunking_params: Dict[str, Any] = Field(
```

```
    default_factory=dict,
```

```
    description="Parameters: chunk_size, overlap, etc."
```

```
)
```

```

processing_timestamp: datetime = Field(
    default_factory=datetime.utcnow,
    description="When chunk was processed"
)

# ===== LIFECYCLE =====
deprecated: bool = Field(default=False, description="Is document/chunk deprecated?")
deprecated_date: Optional[datetime] = Field(None, description="When deprecated")
deprecation_reason: Optional[str] = Field(None, description="Why deprecated")
superseded_by: Optional[str] = Field(None, description="doc_id that replaces this")

# ===== QUALITY & AUTHORITY =====
confidence_score: Optional[float] = Field(
    None,
    ge=0.0,
    le=1.0,
    description="Confidence score for this chunk"
)
authority_level: AuthorityLevel = Field(
    default=AuthorityLevel.DRAFT,
    description="Authority level of document"
)
review_status: ReviewStatus = Field(
    default=ReviewStatus.PENDING,
    description="Review/approval status"
)
reviewed_by: Optional[str] = Field(None, description="Reviewer user ID")
reviewed_date: Optional[datetime] = Field(None, description="When reviewed")

# ===== CUSTOM/EXTENSIBLE =====
custom_metadata: Dict[str, Any] = Field(
    default_factory=dict,
    description="Domain-specific custom fields"
)

@validator('source_file_hash')
def validate_hash(cls, v):
    """Ensure hash is valid SHA-256"""
    if not all(c in '0123456789abcdef' for c in v.lower()):
        raise ValueError('Invalid SHA-256 hash format')
    return v.lower()

class Config:
    use_enum_values = True
    schema_extra = {
        "example": {
            "doc_id": "HR-POLICY-2025-001",
            "chunk_id": "HR-POLICY-2025-001_chunk_5",
            "chunk_text": "Employees are entitled to 15 days of annual leave...",
            "title": "Employee Leave Policy 2025",
            "page_num": 3,
            "char_range": (450, 850),
            "domain": "hr",
            "doc_type": "policy",
            "tags": ["leave", "benefits", "2025"],
            "author": "Jane Doe, HR Director",
        }
    }

```

```

    "uploader_id": "user123",
    "upload_timestamp": "2025-11-20T10:30:00Z",
    "source_file": "Leave_Policy_2025.pdf",
    "source_file_hash": "a3b2c1d4e5f6...",
    "version": "2.1",
    "embedding_model_name": "all-MiniLM-L6-v2",
    "embedding_dimension": 384,
    "chunking_strategy": "recursive",
    "chunking_params": {"chunk_size": 500, "overlap": 50},
    "deprecated": False,
    "authority_level": "official",
    "review_status": "approved"
}
}

```

## 4.2 Required Fields by Operation

Operation	Required Metadata Fields
<b>Initial Upload</b>	doc_id, title, domain, doc_type, uploader_id, source_file, source_file_hash
<b>Chunk Storage</b>	All upload fields + chunk_id, chunk_text, embedding_version, embedding_model_name, embedding_dimension, chunking_strategy, chunking_params
<b>Retrieval Filter</b>	domain, deprecated, doc_type, authority_level, review_status
<b>Deprecation</b>	doc_id, deprecated=True, deprecated_date, deprecation_reason
<b>Version Update</b>	doc_id, version, document_version, last_updated_timestamp, previous_version_id

## 4.3 Metadata Enrichment Pipeline

### During Upload:

1. User provides: doc\_id, title, domain, doc\_type, author
2. System auto-generates: upload\_timestamp, source\_file\_hash
3. System extracts: page\_num, char\_range (from parser)

### During Processing:

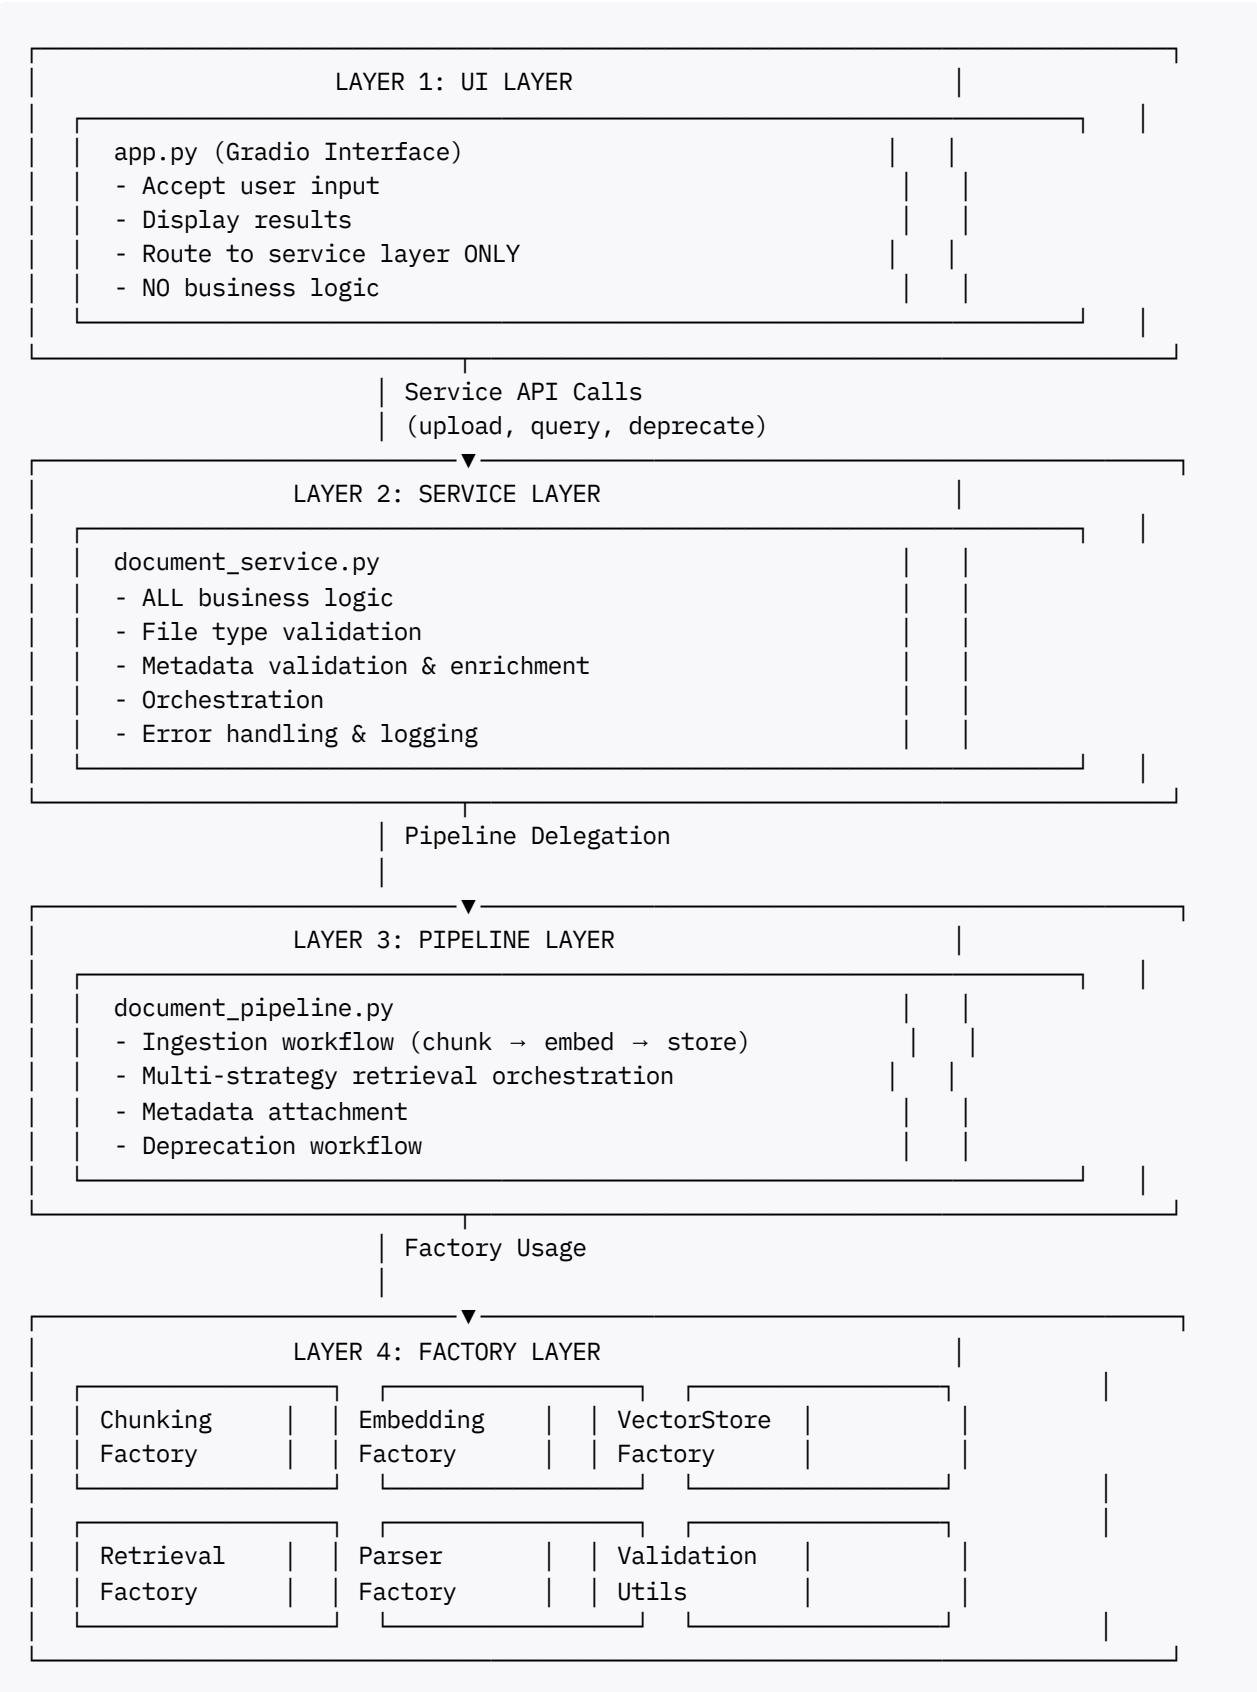
1. System adds: embedding\_version, embedding\_model\_name, embedding\_dimension
2. System adds: chunking\_strategy, chunking\_params, processing\_timestamp
3. System sets defaults: deprecated=False, authority\_level=draft, review\_status=pending

### During Lifecycle Events:

1. Deprecation: Sets deprecated=True, deprecated\_date, deprecation\_reason
2. Update: Increments version, sets last\_updated\_timestamp, links previous\_version\_id
3. Review: Updates review\_status, reviewed\_by, reviewed\_date

## 5. Architecture Overview

### 5.1 Four-Layer Architecture



## 5.2 Data Flow Diagrams

### Ingestion Flow

```
User uploads file via UI
    ↓
UI: Capture file + metadata
    ↓
UI → Service: document_service.upload_document(file, metadata)
    ↓
Service: Validate file type (allowed_file_types check)
Service: Validate metadata (required fields check)
Service: Compute file hash
    ↓
Service → Pipeline: pipeline.process_document(file, enriched_metadata)
    ↓
Pipeline: Extract text via parser factory
Pipeline: Chunk via chunking factory
Pipeline: Embed via embedding factory
Pipeline: Attach metadata to each chunk
Pipeline: Store via vector store factory
    ↓
Pipeline → Service: Return {doc_id, chunks_ingested, status}
    ↓
Service → UI: Return success message
    ↓
UI: Display "Success: 42 chunks ingested"
```

### Query Flow

```
User enters query via UI
    ↓
UI: Capture query text + optional filters
    ↓
UI → Service: document_service.query(query, strategy, filters, top_k)
    ↓
Service: Validate inputs
    ↓
Service → Pipeline: pipeline.query(query, strategy, filters, top_k)
    ↓
Pipeline: Route to retrieval strategy (hybrid, vector, rerank)
Pipeline → Retrieval: Execute search
    ↓
Retrieval: Embed query
Retrieval: Search vector store (dense)
Retrieval: Search BM25 index (sparse) [if hybrid]
Retrieval: Combine scores with alpha weighting [if hybrid]
Retrieval: Apply metadata filters
Retrieval: Return ranked results
    ↓
Pipeline → Service: Return results with metadata
    ↓
Service → UI: Return formatted results
```

↓  
UI: Display results with citations

## 5.3 Component Interaction Matrix

Component	Can Call	Cannot Call	Reason
<b>UI Layer</b>	Service APIs only	Pipeline, Factories, Utils	Separation of concerns
<b>Service Layer</b>	Pipeline, Utils	Factories directly, Vector Stores	Factories via pipeline only
<b>Pipeline Layer</b>	All Factories, Utils	UI layer	Unidirectional flow
<b>Factories</b>	Implementations, Utils	UI, Service, Pipeline	Dependency inversion

## 6. Service Layer Specification

### 6.1 DocumentService Class API

The DocumentService class is the **sole interface** between UI and core business logic.

```
"""
core/services/document_service.py

Service layer for all document operations.
UI MUST ONLY call methods from this class.
"""

import logging
from typing import Dict, Any, Optional, List
from datetime import datetime

from core.pipeline.document_pipeline import DocumentPipeline
from core.config_manager import ConfigManager
from core.utils.validation import validate_file_type, validate_metadata, validate_file_size
from core.metadata_models import ChunkMetadata

logger = logging.getLogger("DocumentService")

class ValidationError(Exception):
    """Raised when validation fails"""
    pass

class DocumentNotFoundError(Exception):
    """Raised when document not found"""
    pass

class DocumentService:
    """
    Service layer providing high-level document management APIs.

    Responsibilities:
    - Input validation (files, metadata)
    """

    def __init__(self, config_manager: ConfigManager):
        self.config_manager = config_manager
        self.pipeline = DocumentPipeline(config_manager)
        self.validator = ValidationHandler(self.config_manager)

    def upload_document(self, file: File, metadata: Dict[str, Any]) -> Document:
        # Validate input
        validate_file_type(file)
        validate_metadata(metadata)

        # Create chunked metadata
        chunks = self.validator.create_chunks(file, metadata)
        chunk_ids = [chunk.id for chunk in chunks]

        # Upload chunks to pipeline
        for chunk in chunks:
            self.pipeline.add_chunk(chunk)

        return Document(id=chunk_ids[0], file=file, metadata=metadata, chunks=chunk_ids)
```

```

- Business logic enforcement
- Pipeline orchestration
- Error handling and logging
- Metadata enrichment

UI Layer MUST call only these methods.
"""

def __init__(self, domain_id: str):
    """
    Initialize service for specific domain.

    Args:
        domain_id: Domain identifier (hr, finance, legal, etc.)
    """
    self.domain_id = domain_id
    self.config_manager = ConfigManager()
    self.domain_config = self.config_manager.load_domain_config(domain_id)
    self.pipeline = DocumentPipeline(self.domain_config)

    # Load security settings
    self.allowed_file_types = set(
        ext.lower() for ext in self.domain_config.security.allowed_file_types
    )
    self.max_file_size_mb = self.domain_config.security.max_file_size_mb

    logger.info(f"DocumentService initialized for domain: {domain_id}")

# ===== PUBLIC API METHODS =====

def upload_document(
    self,
    file_obj: Any,
    metadata: Dict[str, Any],
    replace_existing: bool = False
) -> Dict[str, Any]:
    """
    Upload and process a document end-to-end.

    This method:
    1. Validates file type against allowed_file_types
    2. Validates file size
    3. Validates required metadata fields
    4. Computes file hash for provenance
    5. Enriches metadata with system fields
    6. Delegates to pipeline for processing

    Args:
        file_obj: File object with .name attribute
        metadata: Dict with required fields:
            - doc_id (str)
            - title (str)
            - domain (str)
            - doc_type (str)
            - uploader_id (str)
        replace_existing: If True, deletes existing doc before ingestion
    """

```

```

>Returns:
    Dict with:
        - doc_id (str)
        - chunks_ingested (int)
        - status (str)
        - file_hash (str)

>Raises:
    ValidationError: If validation fails
    ProcessingError: If pipeline processing fails
"""

logger.info(f"Upload request: doc_id={metadata.get('doc_id')}, domain={self.domain_id}")

# Step 1: Validate file type
filename = getattr(file_obj, "name", None)
if not filename:
    raise ValidationError("File object missing 'name' attribute")

validate_file_type(filename, self.allowed_file_types)
logger.debug(f"File type validation passed: {filename}")

# Step 2: Validate file size
file_size = getattr(file_obj, "size", None)
if file_size:
    validate_file_size(file_size, self.max_file_size_mb)
    logger.debug(f"File size validation passed: {file_size} bytes")

# Step 3: Validate required metadata
required_fields = ["doc_id", "title", "domain", "doc_type", "uploader_id"]
validate_metadata(metadata, required_fields)
logger.debug("Metadata validation passed")

# Step 4: Compute file hash
from core.utils.hashing import compute_file_hash
file_hash = compute_file_hash(file_obj)
logger.debug(f"Computed file hash: {file_hash[:16]}...")

# Step 5: Enrich metadata
enriched_metadata = {
    **metadata,
    "source_file": filename,
    "source_file_hash": file_hash,
    "upload_timestamp": datetime.utcnow(),
    "domain": self.domain_id, # Enforce domain
}

# Step 6: Delegate to pipeline
try:
    result = self.pipeline.process_document(
        doc=file_obj,
        metadata=enriched_metadata,
        replace_existing=replace_existing
    )
    result["file_hash"] = file_hash
    logger.info(f"Upload successful: {result}")

```

```

        return result
    except Exception as e:
        logger.exception(f"Pipeline processing failed for doc_id={metadata.get('doc_id')}")
        raise

def deprecate_document(
    self,
    doc_id: str,
    reason: str,
    deprecated_date: Optional[datetime] = None
) -> None:
    """
    Mark a document as deprecated.

    Updates all chunks with:
    - deprecated = True
    - deprecated_date = provided or now
    - depreciation_reason = reason

    Deprecated documents are filtered out of retrieval by default.

    Args:
        doc_id: Document identifier
        reason: Human-readable deprecation reason
        deprecated_date: Optional date; defaults to now

    Raises:
        DocumentNotFoundError: If document doesn't exist
    """
    logger.info(f"Deprecating document: {doc_id}, reason: {reason}")

    if deprecated_date is None:
        deprecated_date = datetime.utcnow()

    try:
        self.pipeline.deprecate_document(doc_id, deprecated_date, reason)
        logger.info(f"Document {doc_id} successfully deprecated")
    except Exception as e:
        logger.exception(f"Deprecation failed for doc_id={doc_id}")
        raise

def update_document_metadata(
    self,
    doc_id: str,
    updates: Dict[str, Any]
) -> None:
    """
    Update metadata fields for a document.

    Allows updating:
    - authority_level
    - review_status
    - tags
    - custom_metadata

    Prevents updating:
    - ...
    """

```

```

- doc_id, chunk_id (immutable)
- upload_timestamp (immutable)
- source_file_hash (immutable)

Args:
    doc_id: Document identifier
    updates: Dict of field: new_value

Raises:
    ValidationError: If trying to update immutable fields
    DocumentNotFoundError: If document doesn't exist
"""

immutable_fields = ["doc_id", "chunk_id", "upload_timestamp", "source_file_hash"]

for field in immutable_fields:
    if field in updates:
        raise ValidationError(f"Cannot update immutable field: {field}")

logger.info(f"Updating metadata for doc_id={doc_id}: {updates}")

try:
    self.pipeline.update_document_metadata(doc_id, updates)
    logger.info(f"Metadata update successful for {doc_id}")
except Exception as e:
    logger.exception(f"Metadata update failed for doc_id={doc_id}")
    raise

def query(
    self,
    query_text: str,
    strategy: Optional[str] = None,
    metadata_filters: Optional[Dict[str, Any]] = None,
    top_k: int = 10,
    include_DEPRECATED: bool = False
) -> List[Dict[str, Any]]:
    """
    Execute semantic query over domain documents.

    Args:
        query_text: Natural language query
        strategy: Retrieval strategy name (hybrid, vector_similarity, etc.)
            If None, uses all configured strategies
        metadata_filters: Optional filters on metadata fields:
            - domain (str or List[str])
            - doc_type (str or List[str])
            - authority_level (str or List[str])
            - tags (str or List[str])
        top_k: Number of results to return
        include_DEPRECATED: If False (default), filters out deprecated docs

    Returns:
        List of dicts with:
            - chunk_id (str)
            - chunk_text (str)
            - score (float)
            - metadata (ChunkMetadata)
    """

```

```

Example:
results = service.query(
    "How many vacation days?",
    strategy="hybrid",
    metadata_filters={"doc_type": "policy", "authority_level": "official"},
    top_k=5
)
"""
logger.info(f"Query: '{query_text}', strategy={strategy}, top_k={top_k}")

# Add default filter for deprecated docs
if not include_DEPRECATED:
    if metadata_filters is None:
        metadata_filters = {}
    metadata_filters["deprecated"] = False

try:
    results = self.pipeline.query(
        query_text=query_text,
        strategy_name=strategy,
        metadata_filters=metadata_filters,
        top_k=top_k
    )
    logger.info(f"Query returned {len(results)} results")
    return results
except Exception as e:
    logger.exception(f"Query failed: {query_text}")
    raise

def get_document_info(self, doc_id: str) -> Dict[str, Any]:
    """
    Retrieve metadata and stats for a document.

    Returns:
        - doc_id
        - title
        - domain
        - upload_timestamp
        - chunk_count
        - deprecated
        - version
        - authority_level
    """
    logger.info(f"Fetching document info: {doc_id}")
    try:
        info = self.pipeline.get_document_info(doc_id)
        return info
    except Exception as e:
        logger.exception(f"Failed to fetch document info: {doc_id}")
        raise DocumentNotFoundError(f"Document not found: {doc_id}")

def list_documents(
    self,
    filters: Optional[Dict[str, Any]] = None,
    include_DEPRECATED: bool = False
)

```

```

) -> List[Dict[str, Any]]:
    """
    List all documents in domain with optional filtering.

    Args:
        filters: Optional metadata filters
        include_deprecated: Include deprecated documents

    Returns:
        List of document summaries
    """
    logger.info(f"Listing documents for domain: {self.domain_id}")

    if not include_deprecated:
        if filters is None:
            filters = {}
        filters["deprecated"] = False

    try:
        docs = self.pipeline.list_documents(filters)
        logger.info(f"Found {len(docs)} documents")
        return docs
    except Exception as e:
        logger.exception("Failed to list documents")
        raise

# ===== VALIDATION HELPERS (used internally) =====

def validate_file_type(self, filename: str) -> bool:
    """Check if file type is allowed. Raises ValidationError if not."""
    return validate_file_type(filename, self.allowed_file_types)

def validate_metadata(self, metadata: Dict, required_fields: List[str]) -> bool:
    """Check if required metadata fields present. Raises ValidationError if not."""
    return validate_metadata(metadata, required_fields)

```

## 6.2 Service Layer Design Principles

Principle	Description
<b>Single Responsibility</b>	Each method does one thing (upload, query, deprecate)
<b>Validation First</b>	All inputs validated before pipeline delegation
<b>Error Transformation</b>	Converts technical errors to user-friendly messages
<b>Logging</b>	Structured logging at every step for observability
<b>No Direct Factory Calls</b>	Always delegates to pipeline; factories via pipeline only
<b>Stateless</b>	No instance state beyond config; thread-safe

## 7. Pipeline Layer Enhancements

### 7.1 Enhanced DocumentPipeline

```
"""
core/pipeline/document_pipeline.py

Enhanced pipeline with:
- Multi-strategy retrieval
- Metadata lifecycle management
- Deprecation support
- Version tracking
"""

import logging
from typing import List, Dict, Optional, Any
from datetime import datetime

from core.factories.chunking_factory import ChunkingFactory
from core.factories.embedding_factory import EmbeddingFactory
from core.factories.vector_store_factory import VectorStoreFactory
from core.factories.retrieval_factory import RetrievalFactory
from core.metadata_models import ChunkMetadata

logger = logging.getLogger("DocumentPipeline")

class DocumentPipeline:
    """
    Orchestrates document processing workflows.

    Responsibilities:
    - Instantiate components via factories
    - Execute chunking → embedding → storage workflow
    - Multi-strategy retrieval orchestration
    - Metadata attachment and lifecycle management
    """

    def __init__(self, domain_config: Any):
        """Initialize pipeline with domain config."""
        self.config = domain_config

        # Create embedding model first (needed for dimension)
        self.embedding_model = EmbeddingFactory.create_embedder(
            self.config.embedding
        )

        # Create chunker
        self.chunker = ChunkingFactory.create_chunker(
            self.config.chunking,
            embedding_model_name=self.embedding_model.model_name
        )

        # Fixed metadata fields
        self.metadata_fields = [
            "doc_id", "chunk_id", "title", "author", "domain",

```

```

        "doc_type", "tags", "upload_timestamp", "version",
        "deprecated", "deprecated_date", "deprecation_reason",
        "source_file", "source_file_hash", "uploader_id",
        "embedding_model_name", "embedding_dimension",
        "chunking_strategy", "chunking_params",
        "authority_level", "review_status"
    ]

    # Create vector store with dimension and metadata schema
    self.vector_store = VectorStoreFactory.create_store(
        self.config.vectorstore,
        embedding_dimension=self.embedding_model.embedding_dimension,
        metadata_fields=self.metadata_fields
    )

    # Initialize retrieval strategies
    self.retrieval_strategies = self._init_retrieval_strategies()

    logger.info(f"DocumentPipeline initialized for domain: {self.config.name}")

def _init_retrieval_strategies(self) -> Dict[str, Any]:
    """Create all configured retrieval strategies."""
    retrieval_cfg = getattr(self.config, "retrieval", {})
    strategies = retrieval_cfg.get("strategies", ["vector_similarity"])

    retrievers = {}
    for strat_name in strategies:
        strat_cfg = retrieval_cfg.get(strat_name, {})
        retriever = RetrievalFactory.create_retriever(
            strat_cfg,
            vector_store=self.vector_store,
            embedding_model=self.embedding_model
        )
        retrievers[strat_name] = retriever
        logger.info(f"Loaded retrieval strategy: {strat_name}")

    return retrievers

def process_document(
    self,
    doc: Any,
    metadata: Dict[str, Any],
    replace_existing: bool = False
) -> Dict[str, Any]:
    """
    Process document end-to-end: chunk → embed → store.

    Args:
        doc: Document file or text
        metadata: Enriched metadata from service layer
        replace_existing: Delete existing before ingestion

    Returns:
        Processing summary
    """
    doc_id = metadata.get("doc_id")

```

```

if not doc_id:
    raise ValueError("metadata must contain 'doc_id'")

logger.info(f"Processing document: {doc_id}")

# Delete existing if requested
if replace_existing:
    self.vector_store.delete_document(doc_id)
    logger.info(f"Deleted existing document: {doc_id}")

# Step 1: Chunk document
chunks = self.chunker.chunk_document(doc)
logger.info(f"Chunked into {len(chunks)} chunks")

# Step 2: Extract text and embed
texts = [c.text for c in chunks]
embeddings = self.embedding_model.embed_texts(texts)
logger.info(f"Generated embeddings for {len(embeddings)} chunks")

# Step 3: Build chunk metadata
chunk_metadatas = []
for i, chunk in enumerate(chunks):
    chunk_meta = {
        # Identity
        "doc_id": doc_id,
        "chunk_id": f"{doc_id}_chunk_{i}",
        "chunk_text": chunk.text,

        # From upload metadata
        "title": metadata.get("title"),
        "author": metadata.get("author"),
        "domain": metadata.get("domain"),
        "doc_type": metadata.get("doc_type"),
        "tags": metadata.get("tags", []),
        "uploader_id": metadata.get("uploader_id"),
        "upload_timestamp": metadata.get("upload_timestamp"),
        "source_file": metadata.get("source_file"),
        "source_file_hash": metadata.get("source_file_hash"),
        "version": metadata.get("version", "1.0"),

        # From chunk
        "page_num": getattr(chunk, "page_num", None),
        "char_range": getattr(chunk, "char_range", None),

        # Processing info
        "embedding_model_name": self.embedding_model.model_name,
        "embedding_dimension": self.embedding_model.embedding_dimension,
        "embedding_version": self.embedding_model.model_name,
        "chunking_strategy": self.config.chunking.strategy,
        "chunking_params": {
            "chunk_size": getattr(self.config.chunking, "chunk_size", None),
            "overlap": getattr(self.config.chunking, "overlap", None)
        },
        "processing_timestamp": datetime.utcnow(),
    }

    # Lifecycle

```

```

        "deprecated": False,
        "deprecated_date": None,
        "deprecation_reason": None,

        # Quality
        "authority_level": metadata.get("authority_level", "draft"),
        "review_status": metadata.get("review_status", "pending"),
    }
    chunk_metadatas.append(chunk_meta)

# Step 4: Upsert to vector store
self.vector_store.upsert_bulk(
    ids=[m["chunk_id"] for m in chunk_metadatas],
    embeddings=embeddings,
    metadatas=chunk_metadatas,
    documents=texts
)
logger.info(f"Upserted {len(chunk_metadatas)} chunks for {doc_id}")

return {
    "doc_id": doc_id,
    "chunks_ingested": len(chunk_metadatas),
    "status": "success"
}

def deprecate_document(
    self,
    doc_id: str,
    deprecated_date: datetime,
    reason: str
) -> None:
    """Mark all chunks for doc_id as deprecated."""
    logger.info(f"Deprecating document: {doc_id}")

    self.vector_store.update_metadata(
        filter={"doc_id": doc_id},
        updates={
            "deprecated": True,
            "deprecated_date": deprecated_date,
            "deprecation_reason": reason
        }
    )
    logger.info(f"Document {doc_id} marked as deprecated")

def update_document_metadata(
    self,
    doc_id: str,
    updates: Dict[str, Any]
) -> None:
    """Update metadata for all chunks of a document."""
    logger.info(f"Updating metadata for {doc_id}: {updates}")

    self.vector_store.update_metadata(
        filter={"doc_id": doc_id},
        updates=updates
)

```

```

        logger.info(f"Metadata updated for {doc_id}")

    def query(
        self,
        query_text: str,
        strategy_name: Optional[str] = None,
        metadata_filters: Optional[Dict[str, Any]] = None,
        top_k: Optional[int] = None
    ) -> List[Dict[str, Any]]:
        """
        Execute query using configured retrieval strategies.

        If strategy_name is None, uses all strategies and merges results.
        """
        if not strategy_name:
            # Use all strategies
            all_results = []
            for name, retriever in self.retrieval_strategies.items():
                results = retriever.retrieve(
                    query_text,
                    metadata_filters=metadata_filters,
                    top_k=top_k
                )
                # Tag with strategy for transparency
                for r in results:
                    r["retrieval_strategy"] = name
                all_results.extend(results)

            # Optional: deduplicate and re-rank
            all_results = self._deduplicate_results(all_results)
            logger.info(f"Multi-strategy query returned {len(all_results)} results")
            return all_results[:top_k] if top_k else all_results
        else:
            # Use specific strategy
            retriever = self.retrieval_strategies.get(strategy_name)
            if not retriever:
                raise ValueError(f"Retrieval strategy '{strategy_name}' not found")

            results = retriever.retrieve(
                query_text,
                metadata_filters=metadata_filters,
                top_k=top_k
            )
            logger.info(f"Query with '{strategy_name}' returned {len(results)} results")
            return results

    def _deduplicate_results(self, results: List[Dict]) -> List[Dict]:
        """Remove duplicate chunks; keep highest score."""
        seen = {}
        for r in results:
            chunk_id = r.get("chunk_id")
            if chunk_id not in seen or r["score"] > seen[chunk_id]["score"]:
                seen[chunk_id] = r
        return sorted(seen.values(), key=lambda x: x["score"], reverse=True)

    def get_document_info(self, doc_id: str) -> Dict[str, Any]:

```

```

    """Fetch document metadata and stats."""
    return self.vector_store.get_document_info(doc_id)

    def list_documents(self, filters: Optional[Dict] = None) -> List[Dict]:
        """List all documents with optional filters."""
        return self.vector_store.list_documents(filters)

```

## 8. Hybrid Retrieval Implementation

### 8.1 Why Hybrid Retrieval?

#### Problem with Pure Vector Search:

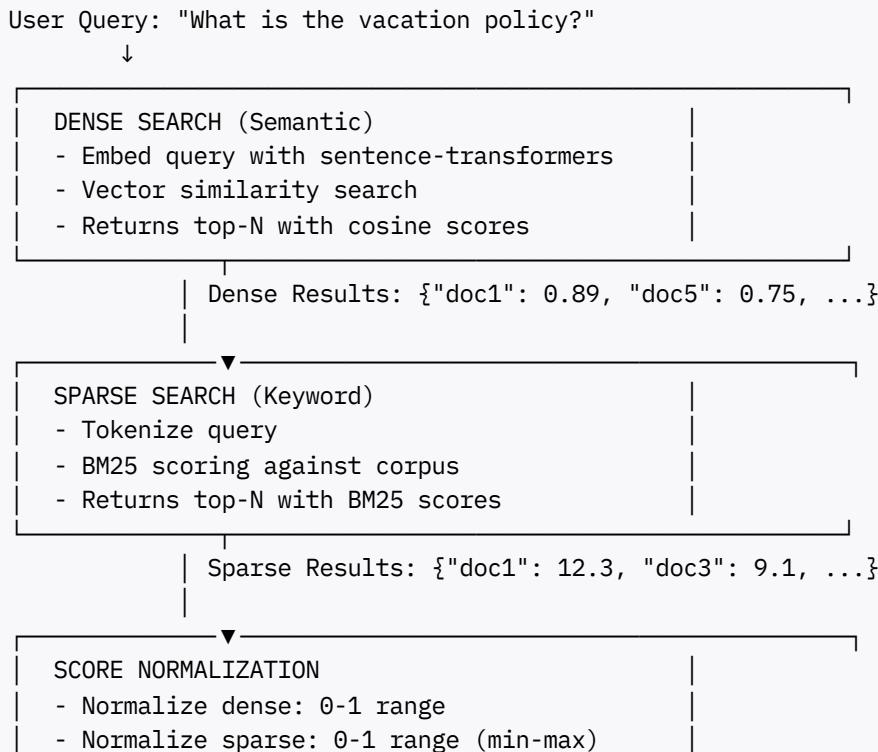
- Misses exact keyword matches
- Poor on domain-specific jargon, acronyms, identifiers
- Less effective on precise terminology (e.g., "Form W-2" vs semantically similar but wrong forms)

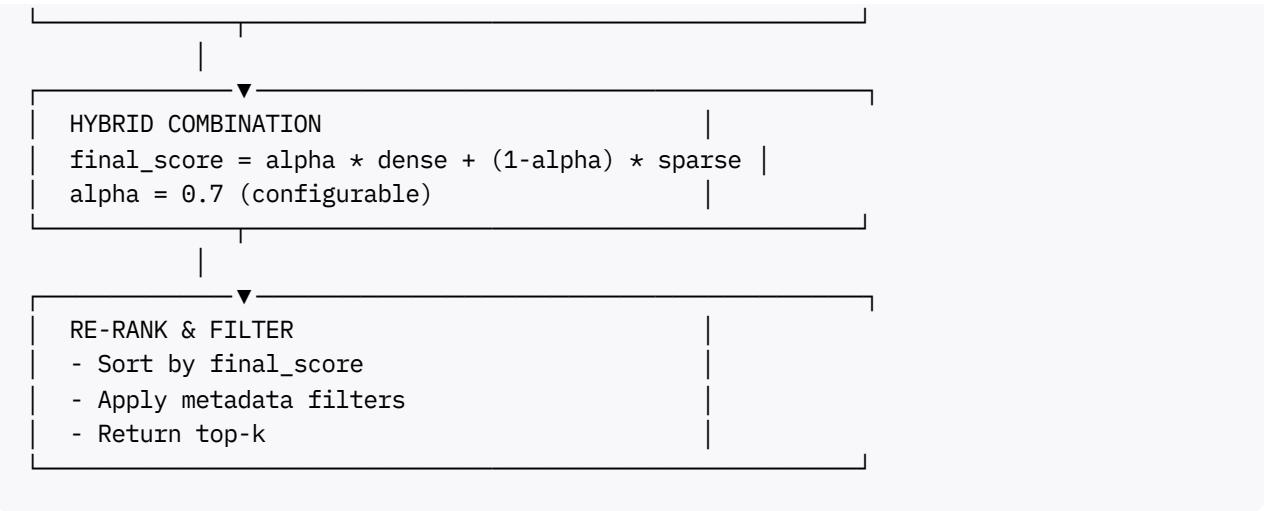
#### Problem with Pure Keyword Search:

- Misses semantic similarity
- Requires exact wording
- No understanding of synonyms or concepts

#### Solution: Hybrid = Dense (Vector) + Sparse (Keyword)

## 8.2 Architecture





### 8.3 BM25 Implementation

```

"""
core/retrievals/bm25_retrieval.py

BM25 sparse keyword retrieval for hybrid search.
"""

import logging
from rank_bm25 import BM25Okapi
from typing import List, Dict, Any
import numpy as np

logger = logging.getLogger("BM25Retrieval")

class BM25Retrieval:
    """
    BM25 (Best Match 25) sparse retrieval.

    Good for:
    - Exact keyword matching
    - Domain-specific terminology
    - Acronyms and identifiers
    - Short queries with specific terms
    """

    def __init__(self, corpus: List[str], doc_ids: List[str]):
        """
        Initialize BM25 index.

        Args:
            corpus: List of document texts
            doc_ids: Corresponding document/chunk IDs
        """
        self.doc_ids = doc_ids
        self.tokenize = lambda text: text.lower().split()
        tokenized_corpus = [self.tokenize(doc) for doc in corpus]
        self.bm25 = BM25Okapi(tokenized_corpus)
        logger.info(f"BM25 index built with {len(corpus)} documents")

```

```

def search(self, query: str, top_k: int = 10) -> List[Dict[str, Any]]:
    """
    Search using BM25.

    Returns:
        List of {doc_id, score} sorted by score descending
    """
    tokenized_query = self.tokenize(query)
    scores = self.bm25.get_scores(tokenized_query)

    # Get top-k indices
    top_indices = np.argsort(scores)[-top_k:][::-1]

    results = [
        {"doc_id": self.doc_ids[i], "score": float(scores[i])}
        for i in top_indices
    ]

    logger.debug(f"BM25 search returned {len(results)} results")
    return results

def normalize_scores(self, scores: List[float]) -> List[float]:
    """Min-max normalization to 0-1 range."""
    if not scores:
        return []
    min_score = min(scores)
    max_score = max(scores)
    if max_score == min_score:
        return [1.0] * len(scores)
    return [(s - min_score) / (max_score - min_score) for s in scores]

```

## 8.4 Hybrid Retrieval Implementation

```

"""
core/retrievals/hybrid_retrieval.py

Combines dense (vector) and sparse (BM25) retrieval with alpha weighting.
"""

import logging
from typing import List, Dict, Any, Optional

logger = logging.getLogger("HybridRetrieval")

class HybridRetrieval:
    """
    Hybrid retrieval combining dense and sparse search.

    Formula:
        final_score = alpha * dense_score + (1 - alpha) * sparse_score

    Where:
        - alpha ∈ [0, 1]: weight parameter
        - alpha = 1.0: pure dense (semantic)
        - alpha = 0.0: pure sparse (keyword)
    """

```

```

    - alpha = 0.7: recommended balanced setting
"""

def __init__(
    self,
    vector_store: Any,
    embedding_model: Any,
    bm25_index: Any,
    alpha: float = 0.7,
    normalize_scores: bool = True
):
    """
    Initialize hybrid retrieval.

    Args:
        vector_store: Dense vector store
        embedding_model: Embedding model for query
        bm25_index: BM25 index for sparse search
        alpha: Weighting parameter (0-1)
        normalize_scores: Whether to normalize before combining
    """
    self.vector_store = vector_store
    self.embedding_model = embedding_model
    self.bm25_index = bm25_index
    self.alpha = alpha
    self.normalize_scores = normalize_scores

    logger.info(f"HybridRetrieval initialized with alpha={alpha}")

def retrieve(
    self,
    query_text: str,
    metadata_filters: Optional[Dict[str, Any]] = None,
    top_k: int = 10
) -> List[Dict[str, Any]]:
    """
    Execute hybrid retrieval.

    Steps:
    1. Dense search (vector similarity)
    2. Sparse search (BM25)
    3. Normalize scores
    4. Combine with alpha weighting
    5. Sort and return top-k
    """
    logger.info(f"Hybrid retrieval: query='{query_text}', alpha={self.alpha}, top_k={top_k}")

    # Step 1: Dense search
    query_embedding = self.embedding_model.embed_texts([query_text])[0]
    dense_results = self.vector_store.search(
        query_embedding,
        top_k=top_k * 2,  # Over-fetch for better recall
        filters=metadata_filters
    )
    dense_dict = {r["chunk_id"]: r["score"] for r in dense_results}
    logger.debug(f"Dense search returned {len(dense_results)} results")

```

```

# Step 2: Sparse search
sparse_results = self.bm25_index.search(query_text, top_k=top_k * 2)
sparse_dict = {r["doc_id"]: r["score"] for r in sparse_results}
logger.debug(f"Sparse search returned {len(sparse_results)} results")

# Step 3: Normalize scores
if self.normalize_scores:
    dense_scores = list(dense_dict.values())
    sparse_scores = list(sparse_dict.values())

    dense_dict = self._normalize_dict(dense_dict)
    sparse_dict = self._normalize_dict(sparse_dict)

# Step 4: Combine scores
all_doc_ids = set(dense_dict.keys()) | set(sparse_dict.keys())
combined = {}

for doc_id in all_doc_ids:
    dense_score = dense_dict.get(doc_id, 0.0)
    sparse_score = sparse_dict.get(doc_id, 0.0)
    combined[doc_id] = self.alpha * dense_score + (1 - self.alpha) * sparse_score

# Step 5: Sort and fetch documents
ranked = sorted(combined.items(), key=lambda x: x[1], reverse=True)[:top_k]

# Fetch full documents with metadata
results = []
for chunk_id, score in ranked:
    doc = self.vector_store.get_document(chunk_id)
    if doc:
        results.append({
            "chunk_id": chunk_id,
            "chunk_text": doc.get("chunk_text"),
            "score": score,
            "metadata": doc.get("metadata"),
            "dense_score": dense_dict.get(chunk_id, 0.0),
            "sparse_score": sparse_dict.get(chunk_id, 0.0)
        })

logger.info(f"Hybrid retrieval returned {len(results)} results")
return results

def _normalize_dict(self, score_dict: Dict[str, float]) -> Dict[str, float]:
    """Min-max normalize scores to 0-1 range."""
    scores = list(score_dict.values())
    if not scores:
        return {}

    min_score = min(scores)
    max_score = max(scores)

    if max_score == min_score:
        return {k: 1.0 for k in score_dict}

    return {

```

```

        k: (v - min_score) / (max_score - min_score)
    for k, v in score_dict.items()
}

```

## 8.5 Alpha Tuning Guidelines

Alpha Value	Behavior	Use Case
<b>1.0</b>	Pure semantic (dense only)	Conceptual questions, paraphrased queries
<b>0.9</b>	Heavy semantic, light keyword	General knowledge questions
<b>0.7</b>	<b>Recommended balanced</b>	Most production use cases
<b>0.5</b>	Equal dense + sparse	Hybrid queries with technical terms
<b>0.3</b>	Keyword-heavy	Specific terminology, codes, identifiers
<b>0.0</b>	Pure keyword (sparse only)	Exact phrase matching, acronyms

### Tuning Process:

1. Start with alpha=0.7
2. Run Golden QA sets
3. Measure Recall@K and MRR
4. Adjust alpha based on domain:
  - Legal/Compliance → lower alpha (more keyword)
  - HR/General knowledge → higher alpha (more semantic)
5. A/B test in production

## 9. Factory Layer Enhancements

### 9.1 Retrieval Factory with Hybrid Support

```

"""
core/factories/retrieval_factory.py

Enhanced to support hybrid retrieval.
"""

import logging
from typing import Any

from core.retrievals.vector_similarity_retrieval import VectorSimilarityRetrieval
from core.retrievals.hybrid_retrieval import HybridRetrieval
from core.retrievals.bm25_retrieval import BM25Retrieval

logger = logging.getLogger("RetrievalFactory")

class RetrievalFactory:

```

```

"""Create retrieval strategies from config."""

_available_strategies = {
    "vector_similarity": VectorSimilarityRetrieval,
    "hybrid": HybridRetrieval,
    # Add more as needed
}

@staticmethod
def create_retriever(
    config: Any,
    vector_store: Any,
    embedding_model: Any,
    bm25_index: Any = None
) -> Any:
    """
    Create retriever based on config.

    For hybrid strategy, requires bm25_index.
    """

    strategy = getattr(config, "strategy", "vector_similarity")
    strategy = strategy.lower()

    logger.info(f"Creating retriever: {strategy}")

    retriever_cls = RetrievalFactory._available_strategies.get(strategy)
    if not retriever_cls:
        raise ValueError(f"Unknown retrieval strategy: {strategy}")

    if strategy == "hybrid":
        if not bm25_index:
            # Build BM25 index from vector store corpus
            corpus, doc_ids = vector_store.get_all_documents()
            bm25_index = BM25Retrieval(corpus, doc_ids)

        alpha = getattr(config, "alpha", 0.7)
        normalize = getattr(config, "normalize_scores", True)

        return HybridRetrieval(
            vector_store=vector_store,
            embedding_model=embedding_model,
            bm25_index=bm25_index,
            alpha=alpha,
            normalize_scores=normalize
        )
    else:
        # Vector similarity or others
        top_k = getattr(config, "top_k", 10)
        return retriever_cls(
            vector_store=vector_store,
            embedding_model=embedding_model,
            top_k=top_k
        )

```

## 10. File Processing & Validation

### 10.1 Validation Utilities

```
"""
core/utils/validation.py

Validation functions for service layer.
"""

from typing import List, Set, Dict, Any

class ValidationError(Exception):
    """Raised when validation fails."""
    pass

def validate_file_type(filename: str, allowed_types: Set[str]) -> bool:
    """
    Check if file extension is allowed.

    Args:
        filename: File name with extension
        allowed_types: Set of allowed extensions (lowercase)

    Raises:
        ValidationError: If file type not allowed

    Returns:
        True if valid
    """
    ext = filename.rsplit('.', 1)[-1].lower() if '.' in filename else ''

    if ext not in allowed_types:
        raise ValidationError(
            f"File type '{ext}' not allowed. "
            f"Allowed types: {', '.join(sorted(allowed_types))}"
        )
    return True

def validate_metadata(metadata: Dict[str, Any], required_fields: List[str]) -> bool:
    """
    Check if required metadata fields are present and non-empty.

    Args:
        metadata: Metadata dict
        required_fields: List of required field names

    Raises:
        ValidationError: If any required field missing or empty

    Returns:
        True if valid
    """
    missing = []
    empty = []

    for field in required_fields:
        if field not in metadata or not metadata[field]:
            if not metadata[field]:
                empty.append(field)
            else:
                missing.append(field)
```

```

for field in required_fields:
    if field not in metadata:
        missing.append(field)
    elif not metadata[field]: # None, empty string, empty list
        empty.append(field)

if missing:
    raise ValidationError(f"Missing required metadata fields: {', '.join(missing)}")
if empty:
    raise ValidationError(f"Empty required metadata fields: {', '.join(empty)}")

return True

def validate_file_size(file_size: int, max_size_mb: int) -> bool:
    """
    Check if file size is within limit.

    Args:
        file_size: File size in bytes
        max_size_mb: Maximum allowed size in MB

    Raises:
        ValidationError: If file too large

    Returns:
        True if valid
    """
    max_bytes = max_size_mb * 1024 * 1024

    if file_size > max_bytes:
        raise ValidationError(
            f"File size {file_size / (1024*1024):.2f} MB exceeds "
            f"maximum allowed size of {max_size_mb} MB"
        )

    return True

```

## 10.2 File Hashing Utility

```

"""
core/utils/hashing.py

File hash computation for provenance tracking.
"""

import hashlib
import logging

logger = logging.getLogger("FileHashing")

def compute_file_hash(file_obj: Any, algorithm: str = "sha256") -> str:
    """
    Compute cryptographic hash of file for integrity verification.

```

```

Args:
    file_obj: File object with .read() method
    algorithm: Hash algorithm (sha256, sha1, md5)

Returns:
    Hexadecimal hash string
"""

if algorithm == "sha256":
    hash_obj = hashlib.sha256()
elif algorithm == "sha1":
    hash_obj = hashlib.sha1()
elif algorithm == "md5":
    hash_obj = hashlib.md5()
else:
    raise ValueError(f"Unsupported hash algorithm: {algorithm}")

# Reset file pointer if possible
if hasattr(file_obj, 'seek'):
    file_obj.seek(0)

# Read in chunks for large files
for chunk in iter(lambda: file_obj.read(4096), b ""):
    hash_obj.update(chunk)

# Reset file pointer again for subsequent reads
if hasattr(file_obj, 'seek'):
    file_obj.seek(0)

hash_hex = hash_obj.hexdigest()
logger.debug(f"Computed {algorithm} hash: {hash_hex[:16]}...")

return hash_hex

```

## 10.3 Enhanced PDF Processor

```

"""
utils/fileparsers/pdf_processor.py

Enhanced PDF processor with metadata extraction.
"""

import logging
from typing import Dict, Any, List
import PyPDF2

from core.utils.hashing import compute_file_hash

logger = logging.getLogger("PDFProcessor")

class PDFProcessor:
    """Process PDF files and extract text with metadata."""

    def extract(self, file_path: str) -> Dict[str, Any]:
        """
        Extract text and metadata from PDF.

```

```

>Returns:
{
    "text": str (full text),
    "metadata": {
        "page_count": int,
        "file_hash": str,
        "pages": List[dict] per-page info
    }
}
"""

logger.info(f"Processing PDF: {file_path}")

try:
    with open(file_path, 'rb') as f:
        reader = PyPDF2.PdfReader(f)

        # Extract text and track page info
        full_text = []
        pages = []
        char_offset = 0

        for page_num, page in enumerate(reader.pages, start=1):
            page_text = page.extract_text()
            full_text.append(page_text)

            pages.append({
                "page_num": page_num,
                "char_range": (char_offset, char_offset + len(page_text)),
                "text_length": len(page_text)
            })
            char_offset += len(page_text)

        # Compute file hash
        f.seek(0)
        file_hash = compute_file_hash(f)

        result = {
            "text": "\n".join(full_text),
            "metadata": {
                "page_count": len(reader.pages),
                "file_hash": file_hash,
                "pages": pages
            }
        }

        logger.info(f"PDF processed: {len(reader.pages)} pages, {len(result['text'])}")
        return result

except Exception as e:
    logger.exception(f"Failed to process PDF: {file_path}")
    raise

```

## 11. Implementation Roadmap

### Phase 2.1: Foundation (Week 1)

#### Task 1.1: Enhanced Metadata Models

**File:** core/metadata\_models.py

##### **Deliverables:**

- Complete Pydantic model with all Phase 2 fields
- Validation methods
- Example schemas

##### **Acceptance Criteria:**

- [ ] All metadata fields defined with types
- [ ] Pydantic validation working
- [ ] Unit tests for validation pass
- [ ] Documentation with examples

**Estimated Time:** 2 days

#### Task 1.2: Create Service Layer

**File:** core/services/document\_service.py

##### **Deliverables:**

- Complete DocumentService class
- All public API methods
- File type and metadata validation
- Error handling

##### **Acceptance Criteria:**

- [ ] Service instantiates per domain
- [ ] upload\_document() functional
- [ ] query() functional
- [ ] deprecate\_document() functional
- [ ] Unit tests cover all methods
- [ ] Zero business logic remains in UI

**Estimated Time:** 3 days

## **Task 1.3: Enhance Pipeline**

**File:** core/pipeline/document\_pipeline.py

### **Deliverables:**

- Multi-strategy retrieval support
- Enhanced metadata attachment
- Deprecation workflow
- Improved logging

### **Acceptance Criteria:**

- [ ] Supports multiple retrieval strategies
- [ ] Metadata properly attached to chunks
- [ ] Deprecation API works
- [ ] Integration tests pass

**Estimated Time:** 3 days

## **Phase 2.2: Hybrid Retrieval (Week 2)**

### **Task 2.1: Implement BM25 Retrieval**

**File:** core/retrievals/bm25\_retrieval.py

### **Deliverables:**

- BM25Retrieval class
- Indexing and search methods
- Score normalization

### **Acceptance Criteria:**

- [ ] BM25 index builds from corpus
- [ ] Search returns ranked results
- [ ] Scores normalized correctly
- [ ] Unit tests pass

**Estimated Time:** 2 days

## **Task 2.2: Implement Hybrid Retrieval**

**File:** core/retrievals/hybrid\_retrieval.py

**Deliverables:**

- HybridRetrieval class
- Score combination logic
- Configurable alpha parameter

**Acceptance Criteria:**

- [ ] Combines dense and sparse correctly
- [ ] Alpha parameter adjustable
- [ ] Results properly ranked
- [ ] Integration test shows improvement over pure vector

**Estimated Time:** 2 days

## **Task 2.3: Update Retrieval Factory**

**File:** core/factories/retrieval\_factory.py

**Deliverables:**

- Hybrid strategy registration
- BM25 instantiation logic

**Acceptance Criteria:**

- [ ] Factory creates hybrid retriever from config
- [ ] All parameters passed correctly
- [ ] Unit tests pass

**Estimated Time:** 1 day

## **Task 2.4: Alpha Tuning & Evaluation**

**Deliverables:**

- Test hybrid with alpha values: 0.3, 0.5, 0.7, 0.9
- Golden QA evaluation
- Optimal alpha per domain

**Acceptance Criteria:**

- [ ] Hybrid outperforms pure vector by  $\geq 15\%$
- [ ] Optimal alpha identified per domain

- [ ] Documentation updated with recommendations

**Estimated Time:** 2 days

## Phase 2.3: File Processing & Validation (Week 3)

### Task 3.1: Validation Utilities

**File:** core/utils/validation.py, core/utils/hashing.py

**Deliverables:**

- validate\_file\_type()
- validate\_metadata()
- validate\_file\_size()
- compute\_file\_hash()

**Acceptance Criteria:**

- [ ] All validation functions working
- [ ] Clear error messages
- [ ] Unit tests cover edge cases
- [ ] Used by service layer

**Estimated Time:** 1 day

### Task 3.2: Enhance File Parsers

**Files:**

- utils/fileparsers/pdf\_processor.py
- utils/fileparsers/docx\_processor.py
- utils/fileparsers/txt\_processor.py

**Deliverables:**

- Extract page numbers
- Extract character ranges
- Compute file hashes
- Better error handling

**Acceptance Criteria:**

- [ ] Page metadata extracted
- [ ] File hashes computed
- [ ] Errors handled gracefully

- [ ] Integration tests pass

**Estimated Time:** 2 days

### Task 3.3: Refactor UI Layer

**File:** app.py

**Deliverables:**

- Remove all business logic
- Replace with service API calls
- Thin handlers (<20 lines each)

**Acceptance Criteria:**

- [ ] Zero business logic in `app.py`
- [ ] All handlers call service methods
- [ ] No direct factory or pipeline imports
- [ ] Code review passes

**Estimated Time:** 2 days

## Phase 2.4: Testing & Quality (Week 4)

### Task 4.1: Unit Tests

**Directory:** tests/unit/

**Deliverables:**

- Factory tests
- Service layer tests
- Pipeline tests
- Retrieval tests
- Validation tests

**Acceptance Criteria:**

- [ ] 80%+ code coverage
- [ ] All critical paths tested
- [ ] CI/CD integration
- [ ] Test documentation

**Estimated Time:** 3 days

## **Task 4.2: Integration Tests**

**Directory:** tests/integration/

**Deliverables:**

- End-to-end ingestion test
- End-to-end query test
- Multi-strategy retrieval test
- Deprecation workflow test

**Acceptance Criteria:**

- [ ] All workflows tested end-to-end
- [ ] Tests use real config files
- [ ] Cleanup after each test
- [ ] CI/CD integration

**Estimated Time:** 2 days

## **Task 4.3: Golden QA Sets**

**Directory:** tests/golden\_qa/

**Deliverables:**

- Golden QA sets for 2+ domains
- Evaluation script
- Baseline metrics
- Comparison: pure vector vs hybrid

**Acceptance Criteria:**

- [ ] QA sets created with expected answers
- [ ] Automated evaluation working
- [ ] Metrics logged: Recall@K, MRR
- [ ] Hybrid shows improvement

**Estimated Time:** 2 days

## **12. Configuration Management**

## 12.1 Global Config (Enhanced)

```
# configs/global_config.yaml
name: global_default
description: Global defaults for multi-domain RAG system (Phase 2)

chunking:
  strategy: recursive
  chunk_size: 500
  overlap: 50

embedding:
  provider: sentence_transformers
  model_name: all-MiniLM-L6-v2
  device: cpu
  batch_size: 32
  normalize_embeddings: true

retrieval:
  strategies:
    - hybrid
    - filtering
  hybrid:
    alpha: 0.7
    dense_provider: sentence_transformers
    sparse_provider: bm25
    normalize_scores: true
  filtering:
    fields: ["domain", "deprecated", "doc_type", "authority_level"]
    default_filters:
      deprecated: false
    top_k: 10

vectorstore:
  provider: chromadb
  index_type: hnsw
  collection_name: default_collection
  persist_directory: ./data/chromadb

metadata:
  track_versions: true
  enable_deprecation: true
  compute_file_hash: true
  extract_page_numbers: true
  required_fields:
    - doc_id
    - title
    - domain
    - doc_type
    - uploader_id
    - upload_date

security:
  allowed_file_types:
    - pdf
    - docx
```

```

    - txt
max_file_size_mb: 20
require_authentication: false

logging:
  level: INFO
  format: "%(asctime)s - %(name)s - %(levelname)s - %(message)s"
  log_file: ./logs/rag_system.log
  enable_structured_logging: true

```

## 12.2 Domain Config Example

```

# configs/domains/hr_domain.yaml
name: hr_domain
description: HR policies and procedures

extends: global_config

retrieval:
  strategies:
    - hybrid
  hybrid:
    alpha: 0.8 # More semantic for HR queries
    top_k: 10

metadata:
  required_fields:
    - doc_id
    - title
    - domain
    - doc_type
    - uploader_id
    - upload_date
    - authority_level

```

## 13. Testing Strategy

### 13.1 Unit Testing

**Coverage Target:** ≥ 80%

**Test Structure:**

```

tests/unit/
├── test_services/
│   └── test_document_service.py
├── test_pipeline/
│   └── test_document_pipeline.py
└── test_factories/
    ├── test_chunking_factory.py
    └── test_embedding_factory.py

```

```

    └── test_retrieval_factory.py
    └── test_vector_store_factory.py
└── test_retrievals/
    ├── test_bm25_retrieval.py
    └── test_hybrid_retrieval.py
└── test_utils/
    ├── test_validation.py
    └── test_hashing.py

```

### Example Unit Test:

```

# tests/unit/test_services/test_document_service.py

import pytest
from core.services.document_service import DocumentService, ValidationError

def test_upload_document_validates_file_type():
    """Service should reject disallowed file types."""
    service = DocumentService("hr")

    # Mock file with disallowed extension
    class MockFile:
        name = "doc.exe"

        metadata = {
            "doc_id": "test123",
            "title": "Test",
            "domain": "hr",
            "doc_type": "policy",
            "uploader_id": "user1"
        }

    with pytest.raises(ValidationError, match="not allowed"):
        service.upload_document(MockFile(), metadata)

def test_query_filters_deprecated_by_default():
    """Service should filter deprecated docs unless explicitly requested."""
    service = DocumentService("hr")

    # Mock pipeline query
    service.pipeline.query = lambda *args, **kwargs: kwargs.get("metadata_filters", {})

    result = service.query("test query")

    # Should have deprecated=False in filters
    assert result.get("deprecated") == False

```

## 13.2 Integration Testing

### Test Scenarios:

1. End-to-end ingestion
2. End-to-end query

3. Multi-strategy retrieval
4. Document lifecycle (upload → query → deprecate → verify filtered)
5. Metadata update workflow

#### **Example Integration Test:**

```
# tests/integration/test_end_to_end.py

import pytest
from core.services.document_service import DocumentService

def test_full_ingestion_query_workflow():
    """Test complete workflow from upload to query."""
    service = DocumentService("test_domain")

    # Upload document
    with open("tests/fixtures/sample.pdf", "rb") as f:
        metadata = {
            "doc_id": "test_doc_001",
            "title": "Test Document",
            "domain": "test_domain",
            "doc_type": "policy",
            "uploader_id": "test_user"
        }
        result = service.upload_document(f, metadata)

    assert result["status"] == "success"
    assert result["chunks_ingested"] > 0

    # Query document
    results = service.query("test query", top_k=5)

    assert len(results) > 0
    assert results[0]["metadata"]["doc_id"] == "test_doc_001"

    # Cleanup
    service.vector_store.delete_document("test_doc_001")
```

### **13.3 Golden QA Sets**

#### **Structure:**

```
# tests/golden_qa/hr_domain_qa.yaml
domain: hr
description: Golden QA set for HR policies

questions:
  - id: hr_001
    question: "How many vacation days do employees get?"
    expected_doc_ids: ["HR-POLICY-2025-001"]
    expected_chunks: ["HR-POLICY-2025-001_chunk_5", "HR-POLICY-2025-001_chunk_6"]
    expected_keywords: ["15 days", "annual leave", "vacation"]
    min_recall_at_5: 1.0
```

```

min_mrr: 0.8

- id: hr_002
  question: "What is the sick leave policy?"
  expected_doc_ids: ["HR-POLICY-2025-001"]
  expected_chunks: ["HR-POLICY-2025-001_chunk_12"]
  expected_keywords: ["sick leave", "10 days", "medical certificate"]
  min_recall_at_5: 1.0
  min_mrr: 0.9

```

### Evaluation Script:

```

# tests/golden_qa/evaluate.py

import yaml
from core.services.document_service import DocumentService

def evaluate_golden_qa(qa_file: str, service: DocumentService):
    """Evaluate retrieval against Golden QA set."""
    with open(qa_file) as f:
        qa_data = yaml.safe_load(f)

    results = []

    for qa in qa_data["questions"]:
        query_results = service.query(qa["question"], top_k=10)

        # Compute metrics
        recall_at_5 = compute_recall(
            retrieved=[r["chunk_id"] for r in query_results[:5]],
            expected=qa["expected_chunks"]
        )

        mrr = compute_mrr(
            retrieved=[r["chunk_id"] for r in query_results],
            expected=qa["expected_chunks"]
        )

        passed = (
            recall_at_5 >= qa["min_recall_at_5"] and
            mrr >= qa["min_mrr"]
        )

        results.append({
            "id": qa["id"],
            "question": qa["question"],
            "recall_at_5": recall_at_5,
            "mrr": mrr,
            "passed": passed
        })

    return results

def compute_recall(retrieved: list, expected: list) -> float:
    """Recall = (retrieved ∩ expected) / expected"""

```

```

if not expected:
    return 1.0
intersection = set(retrieved) & set(expected)
return len(intersection) / len(expected)

def compute_mrr(retrieved: list, expected: list) -> float:
    """Mean Reciprocal Rank."""
    for i, chunk_id in enumerate(retrieved, start=1):
        if chunk_id in expected:
            return 1.0 / i
    return 0.0

```

## 14. CLI Tools

### 14.1 CLI Architecture

```

cli/
└── __init__.py
    ├── ingest.py      # Document ingestion
    ├── query.py       # Query execution
    ├── manage.py      # Domain and document management
    └── evaluate.py   # Golden QA evaluation

```

### 14.2 Ingestion CLI

```

# cli/ingest.py

import click
from core.services.document_service import DocumentService

@click.command()
@click.option('--domain', required=True, help='Domain ID')
@click.option('--file', required=True, type=click.Path(exists=True), help='File path')
@click.option('--doc-id', required=True, help='Document ID')
@click.option('--title', required=True, help='Document title')
@click.option('--doc-type', required=True, help='Document type')
@click.option('--uploader-id', required=True, help='Uploader user ID')
@click.option('--replace', is_flag=True, help='Replace if exists')
def ingest(domain, file, doc_id, title, doc_type, uploader_id, replace):
    """Ingest a document into the RAG system."""
    click.echo(f"Ingesting document: {file} into domain: {domain}")

    service = DocumentService(domain)

    metadata = {
        "doc_id": doc_id,
        "title": title,
        "domain": domain,
        "doc_type": doc_type,
        "uploader_id": uploader_id
    }

```

```
with open(file, 'rb')
```